

## Exercise 02 - Vector addition

In this exercise we will consider the implementation of a kernel for carrying out vector addition. The operation is a fundamental data primitive often used in computations. It is a subclass of the widely used scalar multiplication plus vector addition (SAXPY) computed by the combined operation

$$z = ax + y$$

The generalization of this operation is to conceptually replace vectors with matrices of the same size (see "CUDA Programming Guide" for examples).

### Concepts covered

The following concepts are covered in this exercise

- How to allocate and free memory on GPU.
- How to copy data from CPU to GPU.
- How to copy data from GPU to CPU.
- How to invoke GPU kernels.
- How to write a GPU kernel.

### Files needed

The following files will be needed for the exercise

- `VecAdd.cu` (needs to be updated)
- `VecAdd_kernel.cu` (needs to be updated)

The code will not work until all work steps described below has been successfully completed.

### Work steps

You will need to modify the C code supplied in the files.

The host code in `VecAdd.cu` needs to be updated through the steps

1. Allocate arrays for the vectors  $x$ ,  $y$  and  $z$  in device memory.
2. Transfer arrays  $x$ ,  $y$  and  $z$  from host to device.
3. Define the number of threads per block and blocks per grid to be used in the invocation of the kernel `VecAdd_kernel.cu`.
4. Transfer vector  $c$  from device to host memory.
5. Free allocated device memory for vectors  $x$ ,  $y$  and  $z$ .

The device code in `VecAdd_kernel.cu` needs to be updated through the steps

6. Modify the kernel such that per thread it reads values from  $x$  and  $y$  from global device memory and stores the vector addition result for the element in question in  $z$  back in global device memory.

### Compilation

Compile the final code using the included `Makefile` which includes the common parameters for compilation in `../common.mk`. To compile successfully after CUDA has been successfully installed, you will need to make sure that the environment variables `CUDA_INSTALL_PATH` and `CUDA_SDK_DIR` are correct. In case of compilation errors you will need to debug the code until it compiles successfully.

## Execution

Execute your compiled program. If your program executes successfully the output should look like this

```
Vector addition  
PASSED
```